

Principal AI Platform Engineer

Platform Engineering Team

Engagement Full-time contractor, 40 hours per week	Contract 12 months, contract-to-hire
Working arrangement Remote anywhere in India for the first 12 months	Conversion model Permanent Wolters Kluwer India role, subject to performance and business need
Time-zone overlap 6 hours daily with CST, approx. 7:30 PM - 1:30 AM IST	Post-conversion location Pune or Chennai, hybrid 3 days per week
Cloud Azure strongly preferred; AWS also relevant	Core stack Python, Web APIs, SDKs, GenAI, LLMs and RAG
Assessment GenAI Engineering Skill Test	Interview process Technical and behavioural interview

Role Overview

Wolters Kluwer is seeking a hands-on Principal AI Platform Engineer for the Platform Engineering Team. This is the most senior individual-contributor engineering role on the team. You will set the technical direction for the AI platform, design production-grade AI and software services and SDKs that other teams build on, and raise the delivery bar through rigorous AI-assisted engineering.

The ideal candidate is an elite Python engineer with a track record of owning platform-level architecture at scale, deep expertise in Web APIs, SDK design and cloud platforms, and production-level mastery of LLMs, RAG and prompt engineering. You lead through architecture, standards and influence, not headcount.

Engagement and Location

- **Contract-to-hire:** Start as a full-time contractor for 12 months. Wolters Kluwer intends to convert the role to direct permanent employment in India, subject to performance and business need.
- **India only:** Candidates must be based in India and legally authorised to work there.
- **First 12 months:** Fully remote from anywhere in India.
- **On conversion:** Relocation to Pune or Chennai is required, with a hybrid model of 3 days per week in the office.

Your Skills

Elite Python Engineer: You build resilient, scalable systems and APIs and are the person others rely on for hard design decisions.

Platform Architect: You design systems consumed by other engineering teams and reason clearly about compatibility, versioning, blast radius and total cost of ownership.

AI-Assisted Pro: You use tools such as GitHub Copilot, Cursor or Claude Code daily while maintaining code safety, readability and engineering rigour, and you help define how the wider team does the same.

Quality Driven: You automate testing, builds, deployments and checks, and you establish quality gates that scale across teams.

Cloud-Fluent: You are comfortable with Azure and/or AWS across serverless, containers and managed services.

Strong Communicator: You explain complex trade-offs clearly to technical and non-technical stakeholders, including senior leadership.

Responsibilities

- **Own the architecture** of the AI platform end-to-end, including target-state design, technology selection, migration path and technical standards.
- **Architect and implement** secure, scalable Web APIs in Python using frameworks such as FastAPI, including authentication, authorisation, rate limiting and observability.
- **Design, develop, test, version and publish** Python SDKs consumed by multiple internal teams, including semantic versioning, deprecation policy, documentation and practical samples.
- **Set the standard for AI-assisted development** across the team, including tooling, prompt and review practices, safety guardrails and measurement of productivity gains.
- **Build, deploy and operate** applications on Azure and/or AWS, including Azure Functions, App Service and AKS, or AWS Lambda, ECS and EKS, using infrastructure-as-code where appropriate.
- **Develop production AI solutions** using LLMs and RAG, including Azure OpenAI Service, Amazon Bedrock, embeddings and vector stores, with prompt design, evaluation frameworks and guardrails.
- **Define engineering standards** for unit, integration, contract and load testing, code coverage and quality gates, and make them enforceable in CI.
- **Instrument services** with logging, metrics, tracing and dashboards, lead incident reviews and drive systemic fixes.
- **Partner with product, data, security and platform leadership** to shape the roadmap, define requirements, ship features and meet SLAs and SLOs.
- **Mentor senior and lead engineers** on Python, AI-assisted development, performance, security and operational excellence, and act as a senior technical reviewer for significant designs.

Experience and Skills Required

- **12+ years** of professional software development experience, including **8+ years of hands-on production Python**.
- **Principal or Staff-level scope:** ownership of a platform, service family or SDK used by multiple engineering teams, with clear architectural decisions, trade-offs and measured outcomes.
- **6+ years** building and operating Web APIs at scale using FastAPI, Django or Flask, including OpenAPI/Swagger and contract-first design.
- **4+ years** on Azure and/or AWS using services such as Azure Functions, App Service, AKS, AWS Lambda, ECS or EKS, plus infrastructure-as-code such as Bicep, ARM or Terraform. **Azure is strongly preferred.**
- **3+ years** of hands-on production Generative AI experience, including prompt engineering, LLM integration, RAG pipelines and model or response evaluation frameworks you designed.
- **Daily use of AI coding assistants** such as GitHub Copilot, Cursor or Claude Code, with concrete examples of productivity gains, guardrails and rollout to other engineers.
- **4+ years** designing, versioning and distributing SDKs, including packaging, semantic versioning, PyPI or internal registries, and managing breaking changes across consumers.
- **4+ years** with CI/CD using Azure DevOps and/or GitHub Actions, automated testing with pytest, and quality tooling such as ruff, black and mypy.
- Strong enterprise security knowledge across OAuth2/OIDC, JWT and secrets management.
- Observability experience with centralised logging and tracing, including OpenTelemetry, Azure Monitor or CloudWatch, and experience defining SLIs/SLOs.
- **5+ years** across SQL and NoSQL data layers, plus hands-on use of vector databases such as Azure AI Search, OpenSearch or Pinecone.

Preferred Background

- Degree in Computer Science, Engineering or a related technical field. Strong academic backgrounds from IIT, NIT, BITS, IIIT or comparable institutions are of particular interest.
- Significant engineering experience at high-scale product companies or comparable global product engineering organisations is strongly valued. These are preferences rather than absolute gates.

Nice to Have

- Experience in regulated or enterprise B2B information services such as legal, tax, healthcare, finance or compliance.
- Agentic AI or multi-agent orchestration using LangChain, LangGraph or Semantic Kernel, plus MCP-based tool integration in production.
- Docker, Kubernetes or AKS at scale, including multi-region and high-availability design.
- LLM evaluation and observability tooling, including LLM-as-a-judge, prompt-pipeline tracing, and cost or latency optimisation.
- Open-source contributions, patents, published technical writing or conference speaking.

Soft Skills

- Fluent professional English, written and verbal, with confidence communicating with senior leadership.
- Influences without authority and builds consensus across teams that do not report directly to them.
- Comfortable operating with ambiguity and shifting priorities in a fast-moving AI environment.
- Pragmatic about trade-offs between delivery speed, technical debt and long-term maintainability.

Role Logistics

Working pattern	Full-time contractor role, 40 hours per week.
Time-zone overlap	6 hours daily aligned with CST, approximately 7:30 PM - 1:30 AM IST. This is a firm requirement.
Contract term	12 months as a contractor, followed by intended conversion to permanent Wolters Kluwer India employment, subject to performance and business need.
Location	Remote anywhere in India during the contract phase. On conversion, relocate to Pune or Chennai and work hybrid 3 days per week.
Equipment	Bring your own device.
Assessment	GenAI Engineering Skill Test.
Interview process	Technical and behavioural interview.
Expected start	5 October 2026.

Build the AI platform other engineering teams build on.